

Extended Appendix

1. Further Performance Results

In this section, we present the results from experiments showing the performance of 5 federated learning algorithms learning the CIFAR-10 task using the CNN1 model. We evaluated with a 10 client system with varying number of local client epochs and degree of data heterogeneity with the α parameterized LDA algorithm. This is to compare how our algorithm performed under differing settings. We found that while our algorithm is the best performing in a single epoch and in heterogeneous settings, it is outperformed by the more standard federated optimization algorithms. However, for the 3 to 10 epoch settings we are able to outperform the others with our algorithm by using a modified form of the FedAdam optimizer at the server side. This algorithm is applicable to ours “out-of-the-box” in settings where the aggregated global gradient is available to the server, and we found that it is better performing when the desired β_1, β_2 parameters were raised to the power of the number of local epochs clients performed.

Our algorithm combined with the modified FedAdam is not a catch-all, as for 1 or 2 epoch training there is too much influence from the moments. That is, past gradient estimations influence the current updates too much from overuse, derailing the learning process, while for 3-10 epochs there are sufficient gaps in moment application that past gradients are not misdirecting convergence.

2. Datasets

In this section we provide brief descriptions of the datasets that we have used in our experiments and the pre-processing that we have applied.

2.1. Fashion-MNIST

Fashion-MNIST [1] was created as a drop-in replacement for the immensely popular MNIST dataset [2], and has equivalent structure while providing a more complex learning task. The dataset is composed of 70 000 grayscale images with 28×28 pixels, where each image belongs to one of ten classes, in this case denoting the item of clothing the image is made of (e.g. trouser and dress). 60 000 of these images are in the training dataset, and the other 10 000 in the testing dataset.

For pre-processing, we first scaled the images from their byte-per-pixel encoding into a floating point in the range $[0, 1]$ for each pixel by dividing their value by 255. We then converted each resulting sample into a 3D array with shape $28 \times 28 \times 1$.

TABLE 1. PERFORMANCE OF FEDAVG

α	Epochs	Final accuracy \uparrow
0.1	1	70.363% (0.892%)
	2	67.979% (1.904%)
	3	66.623% (0.787%)
	4	64.704% (1.857%)
	5	62.874% (1.552%)
	10	60.727% (2.411%)
	15	59.605% (3.666%)
	20	58.771% (3.063%)
	25	58.864% (2.390%)
	30	58.106% (2.804%)
0.5	1	69.635% (4.393%)
	2	71.558% (1.138%)
	3	71.261% (0.708%)
	4	70.777% (1.576%)
	5	70.753% (1.039%)
	10	70.085% (1.078%)
	15	69.868% (1.104%)
	20	69.201% (0.957%)
	25	69.488% (1.120%)
	30	69.020% (1.220%)
1.0	1	72.910% (0.516%)
	2	72.703% (0.586%)
	3	72.403% (0.308%)
	4	71.932% (0.324%)
	5	72.439% (0.362%)
	10	72.085% (0.451%)
	15	72.202% (0.739%)
	20	71.548% (0.767%)
	25	71.548% (0.675%)
	30	71.731% (0.507%)
10.0	1	73.034% (1.016%)
	2	73.381% (0.405%)
	3	73.007% (0.267%)
	4	73.271% (0.507%)
	5	73.074% (0.986%)
	10	73.180% (0.417%)
	15	72.897% (0.401%)
	20	73.067% (0.173%)
	25	72.680% (0.485%)
	30	72.726% (0.253%)

2.2. CIFAR-10

CIFAR-10 [3] is an object classification dataset composed of 60 000 RGB color images with 32×32 pixels, with 50 000 samples in the training dataset and 10 000 in the testing dataset. It was created as a subset of the tiny images dataset¹, where reliable class labels are ensured for improved learning quality. The dataset has 10 classes, where each sample is labeled with the one it is an image of (e.g. cat, ship, airplane, etc.).

1. The tiny images dataset was formally withdrawn as specified at <https://groups.csail.mit.edu/vision/TinyImages/>

TABLE 2. PERFORMANCE OF FEDAVG WITH ADAM

α	Epochs	Final accuracy \uparrow
0.1	1	29.016% (3.256%)
	2	51.058% (6.337%)
	3	56.657% (5.801%)
	4	56.317% (4.618%)
	5	57.345% (4.786%)
	10	56.370% (5.781%)
	15	56.147% (5.254%)
	20	52.350% (8.666%)
0.5	25	52.891% (7.930%)
	30	53.539% (3.266%)
	1	60.784% (8.650%)
	2	71.775% (1.063%)
	3	72.346% (1.499%)
	4	72.433% (1.533%)
	5	72.092% (1.704%)
	10	71.631% (2.106%)
1.0	15	71.638% (1.388%)
	20	71.782% (1.493%)
	25	71.805% (1.280%)
	30	71.534% (1.579%)
	1	69.601% (3.237%)
	2	73.875% (1.343%)
	3	74.376% (0.998%)
	4	74.189% (0.637%)
10.0	5	74.215% (0.582%)
	10	74.349% (0.808%)
	15	74.269% (0.486%)
	20	73.985% (0.320%)
	25	74.015% (0.559%)
	30	74.493% (0.775%)
	1	75.501% (0.381%)
	2	75.317% (0.658%)
	3	75.187% (0.119%)
	4	75.765% (0.583%)
	5	75.634% (0.399%)
	10	75.758% (0.488%)
	15	75.771% (0.364%)
	20	76.008% (0.993%)
	25	76.075% (0.579%)
	30	75.931% (1.102%)

TABLE 3. PERFORMANCE OF FEDADAM

α	Epochs	Final accuracy \uparrow
0.1	1	54.050% (3.466%)
	2	9.996% (0.017%)
	3	40.792% (26.681%)
	4	54.400% (1.150%)
	5	58.270% (0.511%)
	10	51.112% (4.641%)
	15	54.858% (3.477%)
	20	54.470% (3.750%)
0.5	25	54.734% (3.307%)
	30	50.942% (4.746%)
	1	27.574% (30.455%)
	2	26.659% (28.853%)
	3	43.480% (29.482%)
	4	63.211% (0.972%)
	5	62.440% (2.560%)
	10	62.951% (0.852%)
1.0	15	62.907% (2.274%)
	20	66.042% (1.441%)
	25	64.370% (1.989%)
	30	64.173% (0.824%)
	1	24.820% (25.667%)
	2	25.494% (26.835%)
	3	61.705% (2.327%)
	4	63.438% (1.275%)
10.0	5	62.043% (2.541%)
	10	65.615% (1.649%)
	15	66.173% (1.564%)
	20	65.398% (3.004%)
	25	67.748% (2.031%)
	30	66.246% (3.975%)
	1	38.866% (25.223%)
	2	10.013% (0.006%)
	3	64.026% (0.614%)
	4	64.373% (2.145%)
	5	66.436% (1.809%)
	10	68.079% (1.182%)
	15	68.950% (0.874%)
	20	68.954% (2.209%)
	25	70.409% (0.794%)
	30	70.376% (0.787%)

To pre-process this dataset, we first standardize the data from the byte-per-pixel-channel format to floating points in the range $[0, 1]$ for each pixel-channel by an element-wise division by 255. We then cast each sample into a 3D array with shape $32 \times 32 \times 3$.

2.3. SVHN

SVHN [4] is a digit recognition dataset similar to the MNIST dataset [2], while providing a significantly harder learning task by using images obtained from real-world house numbers with various natural scenes. We use the cropped digits format of the dataset, which is composed of 99289 RGB color images with 32×32 pixels, where training dataset contains 73257 samples and the testing dataset contains 26032 samples. Each sample is labeled with one of 10 classes stating which digit is contained in the respective image.

3. Projected Gradient Descent

Unless stated otherwise, we have used the projected gradient descent algorithm [5] as an adversarial training step each epoch of training in each experiment. For this algorithm, we set the learning rate, η_{pgd} , to 1% of the primary task learning rate, the maximum amount of perturbation, ϵ , to 0.3, and the number of steps, S , to 1. The specific algorithm we have used is demonstrated in Algorithm 1.

TABLE 4. PERFORMANCE OF OUR PROPOSED ALGORITHM

α	Epochs	Final accuracy \uparrow
0.1	1	70.880% (0.236%)
	2	68.069% (0.794%)
	3	63.478% (1.673%)
	4	59.145% (2.144%)
	5	55.946% (2.558%)
	10	43.717% (2.604%)
	15	32.799% (5.444%)
	20	23.170% (6.714%)
	25	19.224% (4.748%)
	30	22.923% (0.981%)
0.5	1	74.192% (0.502%)
	2	73.264% (1.212%)
	3	70.800% (1.619%)
	4	68.129% (1.210%)
	5	65.658% (1.038%)
	10	50.027% (0.866%)
	15	34.635% (3.083%)
	20	24.947% (5.759%)
	25	22.589% (4.772%)
	30	24.793% (7.588%)
1.0	1	75.307% (0.360%)
	2	74.429% (0.135%)
	3	72.035% (0.539%)
	4	70.336% (0.399%)
	5	67.805% (0.986%)
	10	50.177% (0.417%)
	15	41.069% (2.850%)
	20	21.999% (3.115%)
	25	29.704% (6.101%)
	30	22.196% (7.389%)
10.0	1	75.972% (0.536%)
	2	74.866% (0.273%)
	3	73.054% (0.214%)
	4	71.334% (0.777%)
	5	69.294% (0.807%)
	10	49.593% (1.311%)
	15	41.196% (0.994%)
	20	32.278% (2.564%)
	25	21.645% (3.658%)
	30	24.442% (11.040%)

TABLE 5. PERFORMANCE OF OUR PROPOSED ALGORITHM WITH FEDADAM AGGREGATION

α	Epochs	Final accuracy \uparrow
0.1	1	50.200% (2.480%)
	2	58.327% (4.164%)
	3	63.729% (3.235%)
	4	70.887% (0.729%)
	5	69.408% (1.827%)
	10	67.752% (1.421%)
	15	60.263% (0.678%)
	20	50.651% (6.358%)
	25	47.700% (2.178%)
	30	33.420% (3.145%)
0.5	1	49.890% (4.009%)
	2	64.830% (1.455%)
	3	72.419% (1.663%)
	4	73.725% (1.862%)
	5	74.554% (0.992%)
	10	72.010% (0.935%)
	15	58.691% (5.379%)
	20	58.420% (2.796%)
	25	48.655% (7.004%)
	30	29.760% (8.406%)
1.0	1	34.265% (21.087%)
	2	66.373% (1.567%)
	3	72.786% (0.800%)
	4	75.397% (0.286%)
	5	75.270% (0.256%)
	10	74.245% (0.291%)
	15	59.332% (5.201%)
	20	53.105% (10.382%)
	25	45.600% (2.616%)
	30	28.559% (2.342%)
10.0	1	34.943% (21.605%)
	2	68.884% (0.723%)
	3	74.382% (0.316%)
	4	75.848% (0.336%)
	5	75.671% (1.092%)
	10	73.698% (1.535%)
	15	58.086% (5.665%)
	20	48.865% (6.932%)
	25	41.510% (5.488%)
	30	33.407% (5.987%)

4. Samples of Backdoor Data

Algorithm 1 Projected Gradient Descent

- 1: **function** PGD(objective function evaluated with the current model F_θ , maximum amount of perturbation ϵ , learning rate η_{pgd} , number of steps S , minibatch samples X , minibatch labels y)
 - 2: $X_0 \leftarrow X$
 - 3: **for each** step $s \in [S]$ **do**
 - 4: $g_s \leftarrow \nabla_X F_\theta(X, y)$
 - 5: $X_s \leftarrow X + \eta_{pgd} \text{sign}(g_s)$
 - 6: $X_s \leftarrow \text{clip}(X_s, X - \epsilon, X + \epsilon)$
 - 7: **end for**
 - 8: **return** X_S
 - 9: **end function**
-




Figure 1. Samples of the backdoor data. First column shows the ground truth, second column shows the full pixel pattern trigger, and the last three columns show each of the thirds of the trigger used in the distributed attack.

5. Classifier Metric

For our classifier metric we trained the CNN shown in Table 6 of FMNIST, DenseNet-BC-190 for CIFAR-10, and DenseNet121 for SVHN. Each model was trained for 3000 steps with a minibatch size of 512, using the the yogi [6] optimizer with a learning rate 10^{-3} , $\beta_1 = 0.9$, $\beta_2 = 0.999$, and $\varepsilon = 0.001$.

TABLE 6. CNN USED FOR THE FMNIST CLASSIFIER METRIC

Layer Type	Hyperparameters
Convolutional	32 filters, 3×3 kernel, 2×2 strides
Batch Normalization	Depth axis, $\varepsilon = 1.001 \times 10^{-5}$
ReLU	
Convolutional	32 filters, 3×3 kernel, 2×2 strides
Batch Normalization	Depth axis, $\varepsilon = 1.001 \times 10^{-5}$
ReLU	
Max Pooling	3×3 kernel, 2×2 strides
Convolutional	64 filters, 3×3 kernel, 2×2 strides
Batch Normalization	Depth axis, $\varepsilon = 1.001 \times 10^{-5}$
ReLU	
Convolutional	64 filters, 3×3 kernel, 2×2 strides
Batch Normalization	Depth axis, $\varepsilon = 1.001 \times 10^{-5}$
ReLU	
Global Average Pooling	
Dense	$\#Classes$ neurons
Softmax	

6. Full Mitigation Results

TABLE 7. BASE ALGORITHM AGAINST THE CONTINUOUS ATTACK. CERTIFIED ACCURACY AND CERTIFIED ASR WERE ALWAYS 0 - THEY NEVER RECOVER FROM THE ATTACK

Dataset	Model	Final ACC \uparrow	First ASR \downarrow	Final ASR \downarrow
FMNIST	LeNet	88.258% (0.263%)	0.000% (0.000%)	100.000% (0.000%)
	CNN1	92.004% (0.259%)	27.755% (35.594%)	100.000% (0.000%)
	CNN2	93.059% (0.444%)	89.483% (16.670%)	100.000% (0.000%)
CIFAR-10	LeNet	49.346% (0.067%)	9.913% (6.784%)	38.810% (1.666%)
	CNN1	70.052% (1.407%)	20.531% (8.616%)	45.867% (12.873%)
	CNN2	68.977% (0.629%)	53.427% (38.712%)	42.305% (3.018%)
SVHN	LeNet	68.477% (1.009%)	6.498% (1.207%)	58.879% (7.849%)
	CNN1	85.041% (0.904%)	40.890% (18.386%)	74.173% (9.793%)
	CNN2	84.803% (2.291%)	99.818% (0.273%)	76.091% (3.911%)

TABLE 8. BASE ALGORITHM AGAINST THE ONE-SHOT ATTACK

Dataset	Model	Certified ACC \uparrow	Certified ASR \downarrow	Final ACC \uparrow	First ASR \downarrow	Final ASR \downarrow	Recovery rounds \downarrow
FMNIST	LeNet	87.193% (0.170%)	0.000% (0.000%)	88.318% (0.194%)	8.837% (7.460%)	0.000% (0.000%)	0.667 (0.577)
	CNN1	88.261% (1.846%)	0.000% (0.000%)	90.859% (1.195%)	100.000% (0.000%)	0.000% (0.000%)	0.667 (0.577)
	CNN2	0.000% (0.000%)	0.000% (0.000%)	93.296% (0.061%)	99.059% (1.630%)	0.000% (0.000%)	0.667 (1.155)
CIFAR-10	LeNet	48.274% (1.387%)	11.761% (5.296%)	50.945% (0.406%)	70.867% (50.373%)	7.728% (3.452%)	2.667 (1.155)
	CNN1	67.318% (2.076%)	1.949% (1.030%)	73.444% (0.403%)	91.700% (14.376%)	4.167% (0.254%)	1.000 (0.000)
	CNN2	0.000% (0.000%)	0.000% (0.000%)	72.029% (0.223%)	100.000% (0.000%)	3.797% (0.455%)	1.000 (0.000)
SVHN	LeNet	69.350% (2.274%)	6.663% (5.723%)	72.052% (0.278%)	68.535% (51.844%)	3.042% (0.592%)	1.667 (0.577)
	CNN1	80.841% (2.984%)	4.431% (3.971%)	87.080% (3.754%)	100.000% (0.000%)	1.025% (0.348%)	0.333 (0.577)
	CNN2	0.000% (0.000%)	0.000% (0.000%)	89.377% (0.347%)	67.262% (55.634%)	1.290% (0.216%)	2.333 (2.309)

TABLE 9. NOISING AND CLIPPING METHOD AGAINST THE CONTINUOUS ATTACK

Dataset	Model	Certified ACC \uparrow	Certified ASR \downarrow	Final ACC \uparrow	First ASR \downarrow	Final ASR \downarrow	Recovery rounds \downarrow
FMNIST	LeNet	0.000% (0.000%)	0.000% (0.000%)	65.762% (1.593%)	0.000% (0.000%)	51.109% (16.122%)	0.000 (0.000)
	CNN1	0.000% (0.000%)	0.000% (0.000%)	71.938% (0.562%)	0.000% (0.000%)	62.366% (5.918%)	0.000 (0.000)
	CNN2	0.000% (0.000%)	0.000% (0.000%)	87.987% (0.215%)	0.000% (0.000%)	100.000% (0.000%)	0.000 (0.000)
CIFAR-10	LeNet	0.000% (0.000%)	0.000% (0.000%)	33.771% (0.560%)	6.048% (2.024%)	42.137% (2.358%)	0.000 (0.000)
	CNN1	0.000% (0.000%)	0.000% (0.000%)	39.590% (0.878%)	7.897% (3.001%)	46.875% (5.288%)	0.000 (0.000)
	CNN2	0.000% (0.000%)	0.000% (0.000%)	51.739% (0.872%)	5.242% (1.659%)	35.148% (3.894%)	0.000 (0.000)
SVHN	LeNet	0.000% (0.000%)	0.000% (0.000%)	29.191% (1.036%)	0.942% (0.347%)	75.182% (4.032%)	0.000 (0.000)
	CNN1	0.000% (0.000%)	0.000% (0.000%)	37.582% (0.890%)	2.050% (2.185%)	64.038% (11.371%)	0.000 (0.000)
	CNN2	0.000% (0.000%)	0.000% (0.000%)	70.434% (2.440%)	2.364% (0.949%)	46.032% (5.337%)	0.000 (0.000)

TABLE 10. NOISING AND CLIPPING METHOD AGAINST THE ONE-SHOT ATTACK

Dataset	Model	Certified ACC \uparrow	Certified ASR \downarrow	Final ACC \uparrow	First ASR \downarrow	Final ASR \downarrow	Recovery rounds \downarrow
FMNIST	LeNet	70.940% (0.621%)	0.000% (0.000%)	71.391% (0.880%)	0.000% (0.000%)	0.000% (0.000%)	0.000 (0.000)
	CNN1	71.441% (5.340%)	0.000% (0.000%)	78.536% (0.759%)	0.000% (0.000%)	0.000% (0.000%)	0.000 (0.000)
	CNN2	0.000% (0.000%)	0.000% (0.000%)	88.502% (0.121%)	0.000% (0.000%)	0.000% (0.000%)	0.000 (0.000)
CIFAR-10	LeNet	34.145% (0.643%)	13.710% (6.631%)	37.053% (0.143%)	6.216% (2.017%)	5.880% (1.664%)	999.000 (0.000)
	CNN1	34.178% (7.198%)	10.517% (15.102%)	42.421% (0.772%)	8.367% (3.262%)	6.720% (2.433%)	999.000 (0.000)
	CNN2	0.000% (0.000%)	0.000% (0.000%)	54.497% (0.608%)	5.544% (1.067%)	5.612% (0.915%)	999.000 (0.000)
SVHN	LeNet	28.823% (4.557%)	2.960% (3.549%)	36.075% (0.664%)	0.992% (0.455%)	1.438% (0.473%)	333.000 (576.773)
	CNN1	39.023% (4.453%)	26.571% (36.841%)	42.968% (0.297%)	2.133% (2.181%)	3.869% (2.198%)	342.000 (569.139)
	CNN2	0.000% (0.000%)	0.000% (0.000%)	75.269% (1.912%)	2.331% (0.603%)	2.116% (0.862%)	634.667 (504.377)

TABLE 12. ALGORITHM WITH PGD AGAINST THE ONE-SHOT ATTACK

Dataset	Model	ε	Certified ACC \uparrow	Certified ASR \downarrow	Final ACC \uparrow	First ASR \downarrow	Final ASR \downarrow	Recovery \downarrow
FMNIST	LeNet	0.010	87.216% (0.115%)	0.000% (0.000%)	88.425% (0.182%)	8.199% (7.046%)	0.000% (0.000%)	0.667 (0.577)
		0.100	87.216% (0.115%)	0.000% (0.000%)	88.425% (0.182%)	8.199% (7.046%)	0.000% (0.000%)	0.667 (0.577)
		0.500	0.000% (0.000%)	0.000% (0.000%)	88.425% (0.182%)	8.199% (7.046%)	0.000% (0.000%)	0.667 (0.577)
		1.000	0.000% (0.000%)	0.000% (0.000%)	88.425% (0.182%)	8.199% (7.046%)	0.000% (0.000%)	0.667 (0.577)
	CNN1	0.010	87.520% (2.097%)	0.000% (0.000%)	90.612% (1.404%)	100.000% (0.000%)	0.000% (0.000%)	1.000 (0.000)
		0.100	88.408% (1.165%)	0.000% (0.000%)	90.809% (1.144%)	100.000% (0.000%)	0.000% (0.000%)	1.000 (0.000)
		0.500	0.000% (0.000%)	0.000% (0.000%)	91.149% (0.745%)	100.000% (0.000%)	0.000% (0.000%)	1.000 (0.000)
		1.000	0.000% (0.000%)	0.000% (0.000%)	90.939% (1.128%)	100.000% (0.000%)	0.000% (0.000%)	1.000 (0.000)
	CNN2	0.010	88.061% (3.641%)	1.344% (2.072%)	93.199% (0.226%)	99.866% (0.233%)	0.302% (0.363%)	0.667 (1.155)
		0.100	86.285% (2.890%)	0.302% (0.175%)	93.263% (0.163%)	76.579% (40.566%)	0.437% (0.308%)	0.333 (0.577)
		0.500	0.000% (0.000%)	0.000% (0.000%)	93.269% (0.050%)	100.000% (0.000%)	0.941% (0.616%)	0.333 (0.577)
		1.000	0.000% (0.000%)	0.000% (0.000%)	93.259% (0.050%)	100.000% (0.000%)	0.840% (0.857%)	0.333 (0.577)
CIFAR-10	LeNet	0.010	48.017% (1.527%)	11.122% (6.464%)	51.219% (0.546%)	77.151% (39.576%)	6.384% (2.175%)	2.333 (0.577)
		0.100	48.017% (1.527%)	11.122% (6.464%)	51.219% (0.546%)	77.151% (39.576%)	6.384% (2.175%)	2.333 (0.577)
		0.500	0.000% (0.000%)	0.000% (0.000%)	51.219% (0.546%)	77.151% (39.576%)	6.384% (2.175%)	2.333 (0.577)
		1.000	0.000% (0.000%)	0.000% (0.000%)	51.219% (0.546%)	77.151% (39.576%)	6.384% (2.175%)	2.333 (0.577)
	CNN1	0.010	66.249% (2.992%)	1.882% (2.299%)	73.344% (0.344%)	90.726% (16.063%)	4.234% (1.287%)	1.000 (0.000)
		0.100	67.922% (2.673%)	3.831% (1.310%)	73.604% (0.536%)	91.465% (14.783%)	3.763% (0.324%)	1.000 (0.000)
		0.500	0.000% (0.000%)	0.000% (0.000%)	73.865% (0.551%)	90.894% (15.772%)	4.704% (0.455%)	1.000 (0.000)
		1.000	0.000% (0.000%)	0.000% (0.000%)	73.628% (0.321%)	96.136% (6.693%)	3.999% (0.507%)	0.667 (0.577)
	CNN2	0.010	52.534% (3.204%)	15.323% (20.369%)	71.895% (0.113%)	100.000% (0.000%)	3.696% (0.308%)	1.000 (0.000)
		0.100	52.658% (5.028%)	17.272% (24.941%)	72.115% (0.454%)	67.507% (56.280%)	3.528% (0.561%)	2.000 (1.732)
		0.500	0.000% (0.000%)	0.000% (0.000%)	72.082% (0.278%)	100.000% (0.000%)	3.763% (0.937%)	1.000 (0.000)
		1.000	0.000% (0.000%)	0.000% (0.000%)	57.128% (25.604%)	100.000% (0.000%)	3.595% (1.486%)	1.000 (0.000)
SVHN	LeNet	0.010	68.286% (0.317%)	6.068% (2.669%)	71.004% (1.214%)	66.882% (53.887%)	4.117% (0.700%)	2.333 (0.577)
		0.100	68.286% (0.317%)	6.068% (2.669%)	71.004% (1.214%)	66.882% (53.887%)	4.117% (0.700%)	2.333 (0.577)
		0.500	0.000% (0.000%)	0.000% (0.000%)	71.004% (1.214%)	66.882% (53.887%)	4.117% (0.700%)	2.333 (0.577)
		1.000	0.000% (0.000%)	0.000% (0.000%)	71.004% (1.214%)	66.882% (53.887%)	4.117% (0.700%)	2.333 (0.577)
	CNN1	0.010	80.667% (4.435%)	7.870% (6.594%)	87.749% (2.600%)	100.000% (0.000%)	0.810% (0.159%)	0.333 (0.577)
		0.100	80.841% (3.866%)	2.877% (3.716%)	87.403% (2.739%)	100.000% (0.000%)	0.942% (0.406%)	0.333 (0.577)
		0.500	0.000% (0.000%)	0.000% (0.000%)	87.431% (3.043%)	100.000% (0.000%)	0.843% (0.050%)	0.333 (0.577)
		1.000	0.000% (0.000%)	0.000% (0.000%)	87.432% (3.163%)	100.000% (0.000%)	1.042% (0.358%)	0.333 (0.577)
	CNN2	0.010	81.304% (1.472%)	4.828% (4.024%)	89.408% (0.166%)	89.319% (18.500%)	1.190% (0.050%)	1.333 (0.577)
		0.100	80.759% (2.654%)	3.985% (4.969%)	89.581% (0.052%)	100.000% (0.000%)	1.257% (0.273%)	2.000 (1.000)
		0.500	0.000% (0.000%)	0.000% (0.000%)	89.512% (0.150%)	99.950% (0.086%)	0.810% (0.319%)	1.000 (0.000)
		1.000	0.000% (0.000%)	0.000% (0.000%)	89.532% (0.128%)	100.000% (0.000%)	0.876% (0.076%)	1.000 (0.000)

7. Additional Inversion Results

In this section, we present results from experiments analyzing the impact of the representation inversion attack on algorithms and their variants, which have been omitted from the main paper.

In Table 13 we present the results of applying the simple data perturbation technique of random rotation by $[-10, 10]^\circ$ on the inversion and its impact on the general performance of our proposed algorithm. We see that despite its simplicity, this form of perturbation results in better inversion mitigation in the CNN settings when learning with FMNIST and SVHN. However, there is a cost in the accuracy of the models, and it is not always beneficial as seen with the CIFAR-10 setting. These issues can be potentially addressed with a more advanced data perturbation technique.

TABLE 13. OUR ALGORITHM WITH SIMPLE PERTURBATION AGAINST INVERSION

Dataset	Model	Final ACC \uparrow	PSNR \downarrow	SSIM \downarrow
FMNIST	LeNet	85.068% (0.340%)	8.344 (0.224)	0.039 (0.006)
	CNN1	89.158% (0.709%)	8.580 (0.258)	0.050 (0.013)
	CNN2	89.442% (0.609%)	8.484 (0.222)	0.056 (0.008)
CIFAR-10	LeNet	40.599% (0.318%)	11.240 (0.323)	0.163 (0.009)
	CNN1	65.316% (1.334%)	11.237 (0.364)	0.159 (0.007)
	CNN2	64.702% (0.428%)	11.253 (0.495)	0.163 (0.005)
SVHN	LeNet	23.790% (3.636%)	14.038 (1.910)	0.297 (0.052)
	CNN1	84.805% (2.098%)	14.972 (0.489)	0.323 (0.030)
	CNN2	87.136% (1.097%)	14.693 (0.290)	0.322 (0.019)

In Tables 14–22, we have analyzed the impact of the inversion attack on the federated learning algorithm with noising and clipping, as in [7]. We have exhaustively analyzed the effect of the clipping parameter, C , and the noise standard deviation parameter, σ which contribute to the client model updates in the form $\tilde{\mathbf{u}}_i = \min\left(1, \frac{\|\mathbf{u}_i\|}{C}\right) \mathbf{u}_i + N(0, \sigma^2)$. We see that this algorithm only achieves better inversion mitigation than our proposed algorithm in settings that significantly impact the task performance of the model. This further justifies the design of our proposed algorithm as it does not have such a trade-off.

TABLE 14. IMPACT OF INVERSION ON THE NOISING AND CLIPPING FEDAVG ALGORITHM ON THE FMNIST DATASET WITH THE LeNET MODEL

C	σ	Final ACC \uparrow	PSNR \downarrow	SSIM \downarrow
0.100	0.001	79.210% (0.502%)	8.749 (0.297)	0.072 (0.008)
	0.010	60.478% (1.031%)	7.831 (0.102)	0.027 (0.006)
	0.050	13.694% (3.372%)	7.692 (0.053)	0.021 (0.005)
	0.100	12.994% (3.718%)	7.686 (0.061)	0.021 (0.004)
0.500	0.001	83.313% (0.177%)	8.420 (0.158)	0.058 (0.010)
	0.010	74.483% (1.241%)	9.276 (0.258)	0.072 (0.010)
	0.050	16.643% (2.565%)	7.838 (0.240)	0.022 (0.004)
	0.100	13.247% (3.844%)	7.691 (0.065)	0.021 (0.004)
1.000	0.001	83.507% (0.185%)	8.448 (0.166)	0.066 (0.013)
	0.010	76.718% (1.150%)	9.208 (0.307)	0.071 (0.002)
	0.050	23.659% (2.431%)	8.211 (0.843)	0.039 (0.027)
	0.100	13.307% (3.837%)	7.824 (0.225)	0.022 (0.005)
5.000	0.001	83.587% (0.223%)	8.457 (0.169)	0.067 (0.013)
	0.010	76.831% (1.013%)	9.213 (0.231)	0.077 (0.015)
	0.050	28.153% (4.810%)	8.908 (0.085)	0.079 (0.002)
	0.100	14.201% (2.793%)	7.817 (0.261)	0.025 (0.006)
10.000	0.001	83.587% (0.223%)	8.457 (0.169)	0.067 (0.013)
	0.010	76.831% (1.013%)	9.225 (0.229)	0.077 (0.015)
	0.050	26.611% (1.426%)	8.924 (0.099)	0.079 (0.003)
	0.100	13.524% (2.991%)	7.621 (0.190)	0.024 (0.008)

TABLE 15. IMPACT OF INVERSION ON THE NOISING AND CLIPPING FEDAVG ALGORITHM ON THE FMNIST DATASET WITH THE CNN1 MODEL

C	σ	Final ACC \uparrow	PSNR \downarrow	SSIM \downarrow
0.100	0.001	71.774% (4.343%)	8.522 (0.175)	0.052 (0.008)
	0.010	40.462% (6.267%)	8.886 (0.415)	0.052 (0.009)
	0.050	9.524% (0.812%)	8.636 (0.360)	0.050 (0.027)
	0.100	9.224% (1.332%)	8.756 (0.409)	0.050 (0.027)
0.500	0.001	85.935% (1.571%)	8.817 (0.203)	0.071 (0.002)
	0.010	63.294% (2.288%)	9.365 (0.367)	0.070 (0.014)
	0.050	9.578% (0.720%)	8.572 (0.356)	0.043 (0.033)
	0.100	9.211% (1.355%)	8.683 (0.419)	0.050 (0.029)
1.000	0.001	87.110% (0.976%)	9.018 (0.208)	0.070 (0.009)
	0.010	72.178% (2.184%)	9.119 (0.231)	0.056 (0.008)
	0.050	9.494% (0.864%)	8.589 (0.233)	0.037 (0.020)
	0.100	9.284% (1.228%)	8.688 (0.411)	0.048 (0.029)
5.000	0.001	87.010% (1.104%)	9.337 (0.173)	0.082 (0.009)
	0.010	25.444% (25.601%)	8.675 (0.597)	0.052 (0.019)
	0.050	8.523% (2.519%)	8.725 (0.033)	0.055 (0.015)
	0.100	9.411% (1.008%)	8.607 (0.450)	0.039 (0.020)
10.000	0.001	86.920% (1.209%)	9.358 (0.204)	0.082 (0.011)
	0.010	30.815% (36.021%)	8.951 (0.104)	0.059 (0.021)
	0.050	10.168% (0.304%)	8.576 (0.154)	0.048 (0.023)
	0.100	9.548% (0.772%)	8.618 (0.175)	0.041 (0.013)

TABLE 16. IMPACT OF INVERSION ON THE NOISING AND CLIPPING FEDAVG ALGORITHM ON THE FMNIST DATASET WITH THE CNN2 MODEL

C	σ	Final ACC \uparrow	PSNR \downarrow	SSIM \downarrow
0.100	0.001	83.827% (0.493%)	8.931 (0.387)	0.080 (0.013)
	0.010	62.744% (3.437%)	8.935 (0.200)	0.060 (0.004)
	0.050	11.493% (3.033%)	9.239 (0.250)	0.080 (0.010)
	0.100	6.825% (2.808%)	9.193 (0.164)	0.075 (0.015)
0.500	0.001	87.807% (0.322%)	9.063 (0.412)	0.074 (0.002)
	0.010	77.862% (1.354%)	9.013 (0.366)	0.057 (0.011)
	0.050	18.165% (5.889%)	9.090 (0.241)	0.076 (0.007)
	0.100	10.131% (0.579%)	9.288 (0.091)	0.084 (0.011)
1.000	0.001	88.117% (0.267%)	9.240 (0.344)	0.078 (0.007)
	0.010	79.424% (0.864%)	9.031 (0.184)	0.060 (0.013)
	0.050	13.261% (8.701%)	8.851 (0.123)	0.077 (0.010)
	0.100	11.119% (0.978%)	9.160 (0.072)	0.080 (0.010)
5.000	0.001	87.790% (0.237%)	9.477 (0.179)	0.106 (0.018)
	0.010	78.880% (0.912%)	9.199 (0.230)	0.087 (0.023)
	0.050	10.358% (0.779%)	8.825 (0.103)	0.077 (0.010)
	0.100	10.001% (0.012%)	8.822 (0.098)	0.077 (0.009)
10.000	0.001	86.809% (0.236%)	9.536 (0.068)	0.117 (0.014)
	0.010	79.484% (1.422%)	9.460 (0.138)	0.095 (0.011)
	0.050	9.958% (0.079%)	8.815 (0.086)	0.076 (0.008)
	0.100	12.920% (5.070%)	9.008 (0.227)	0.072 (0.012)

TABLE 17. IMPACT OF INVERSION ON THE NOISING AND CLIPPING FEDAVG ALGORITHM ON THE CIFAR-10 DATASET WITH THE LeNet MODEL

C	σ	Final ACC \uparrow	PSNR \downarrow	SSIM \downarrow
0.100	0.001	34.591% (0.656%)	10.572 (0.554)	0.141 (0.011)
	0.010	13.050% (2.193%)	12.005 (0.207)	0.165 (0.003)
	0.050	9.037% (1.009%)	6.972 (0.315)	0.055 (0.015)
	0.100	9.071% (0.998%)	6.971 (0.323)	0.055 (0.015)
0.500	0.001	40.813% (1.406%)	10.481 (0.334)	0.148 (0.003)
	0.010	22.742% (0.214%)	11.887 (0.233)	0.161 (0.002)
	0.050	9.388% (0.426%)	7.040 (0.449)	0.059 (0.021)
	0.100	9.114% (0.942%)	6.978 (0.325)	0.055 (0.015)
1.000	0.001	40.139% (1.229%)	10.595 (0.414)	0.152 (0.001)
	0.010	14.472% (3.357%)	12.136 (0.093)	0.167 (0.006)
	0.050	9.141% (1.467%)	9.802 (1.460)	0.126 (0.028)
	0.100	9.091% (0.940%)	6.965 (0.344)	0.055 (0.016)
5.000	0.001	39.575% (1.674%)	10.704 (0.397)	0.153 (0.001)
	0.010	14.388% (3.350%)	12.085 (0.093)	0.165 (0.003)
	0.050	9.664% (0.334%)	8.622 (2.996)	0.088 (0.073)
	0.100	9.494% (0.454%)	7.069 (0.316)	0.056 (0.017)
10.000	0.001	39.575% (1.674%)	10.702 (0.398)	0.153 (0.001)
	0.010	14.081% (3.360%)	11.932 (0.166)	0.166 (0.003)
	0.050	9.614% (0.855%)	7.485 (0.324)	0.070 (0.023)
	0.100	9.624% (0.454%)	7.001 (0.304)	0.055 (0.017)

TABLE 18. IMPACT OF INVERSION ON THE NOISING AND CLIPPING FEDAVG ALGORITHM ON THE CIFAR-10 DATASET WITH THE CNN1 MODEL

C	δ	Final ACC \uparrow	PSNR \downarrow	SSIM \downarrow
0.100	0.001	39.215% (0.822%)	11.563 (0.419)	0.171 (0.014)
	0.010	10.612% (0.466%)	10.358 (2.570)	0.139 (0.060)
	0.050	8.827% (1.534%)	11.825 (0.385)	0.170 (0.019)
	0.100	8.754% (1.704%)	11.896 (0.372)	0.170 (0.019)
0.500	0.001	54.517% (1.138%)	12.023 (0.413)	0.168 (0.010)
	0.010	24.113% (3.293%)	10.130 (0.600)	0.138 (0.014)
	0.050	8.647% (1.765%)	11.770 (0.441)	0.169 (0.018)
	0.100	8.844% (1.661%)	11.882 (0.333)	0.170 (0.018)
1.000	0.001	55.034% (0.525%)	12.056 (0.149)	0.170 (0.008)
	0.010	14.101% (3.441%)	10.329 (0.958)	0.141 (0.026)
	0.050	8.540% (2.092%)	11.533 (0.422)	0.164 (0.018)
	0.100	8.810% (1.791%)	11.818 (0.381)	0.170 (0.019)
5.000	0.001	53.459% (3.380%)	12.202 (0.585)	0.170 (0.009)
	0.010	10.045% (0.090%)	10.815 (2.425)	0.146 (0.058)
	0.050	8.432% (2.710%)	11.377 (0.444)	0.173 (0.000)
	0.100	8.667% (1.911%)	11.809 (0.122)	0.175 (0.001)
10.000	0.001	54.181% (0.999%)	12.102 (0.228)	0.170 (0.011)
	0.010	13.327% (5.029%)	11.406 (1.214)	0.162 (0.023)
	0.050	8.122% (1.922%)	11.711 (0.360)	0.170 (0.014)
	0.100	9.083% (1.230%)	11.651 (0.351)	0.170 (0.015)

TABLE 19. IMPACT OF INVERSION ON THE NOISING AND CLIPPING FEDAVG ALGORITHM ON THE CIFAR-10 DATASET WITH THE CNN2 MODEL

C	δ	Final ACC \uparrow	PSNR \downarrow	SSIM \downarrow
0.100	0.001	41.593% (0.957%)	10.086 (0.177)	0.139 (0.014)
	0.010	16.994% (1.474%)	11.670 (0.437)	0.162 (0.013)
	0.050	11.156% (1.570%)	12.096 (0.295)	0.166 (0.010)
	0.100	9.938% (0.174%)	12.066 (0.224)	0.166 (0.010)
0.500	0.001	52.412% (1.042%)	11.531 (0.104)	0.155 (0.011)
	0.010	32.026% (1.576%)	11.033 (0.153)	0.152 (0.004)
	0.050	10.305% (0.603%)	12.056 (0.148)	0.165 (0.007)
	0.100	10.045% (0.129%)	12.044 (0.193)	0.166 (0.010)
1.000	0.001	51.858% (1.305%)	11.656 (0.303)	0.157 (0.009)
	0.010	31.278% (0.675%)	11.162 (0.077)	0.155 (0.009)
	0.050	10.085% (0.099%)	12.166 (0.192)	0.167 (0.010)
	0.100	10.335% (0.575%)	12.150 (0.185)	0.167 (0.009)
5.000	0.001	50.370% (0.180%)	11.538 (0.398)	0.158 (0.014)
	0.010	31.285% (1.632%)	11.281 (0.110)	0.153 (0.008)
	0.050	10.115% (0.139%)	11.997 (0.277)	0.165 (0.006)
	0.100	9.988% (0.020%)	12.117 (0.185)	0.167 (0.010)
10.000	0.001	51.565% (1.745%)	11.772 (0.342)	0.164 (0.011)
	0.010	31.642% (1.362%)	11.129 (0.266)	0.155 (0.011)
	0.050	9.901% (0.159%)	11.966 (0.148)	0.164 (0.006)
	0.100	10.001% (0.006%)	12.130 (0.187)	0.167 (0.009)

TABLE 20. IMPACT OF INVERSION ON THE NOISING AND CLIPPING FEDAVG ALGORITHM ON THE SVHN DATASET WITH THE LeNET MODEL

C	δ	Final ACC \uparrow	PSNR \downarrow	SSIM \downarrow
0.100	0.001	29.247% (0.879%)	11.119 (0.505)	0.222 (0.013)
	0.010	18.734% (0.236%)	15.010 (0.144)	0.326 (0.006)
	0.050	9.253% (4.552%)	8.505 (0.795)	0.109 (0.038)
	0.100	9.267% (4.732%)	8.532 (0.804)	0.109 (0.038)
0.500	0.001	44.373% (4.813%)	12.525 (0.767)	0.264 (0.012)
	0.010	19.336% (0.163%)	14.878 (0.376)	0.327 (0.005)
	0.050	9.186% (4.291%)	8.564 (0.861)	0.113 (0.042)
	0.100	9.358% (4.762%)	8.540 (0.820)	0.110 (0.039)
1.000	0.001	31.275% (2.347%)	11.832 (0.280)	0.248 (0.014)
	0.010	19.422% (0.096%)	14.808 (0.328)	0.323 (0.002)
	0.050	10.469% (3.489%)	11.494 (1.861)	0.231 (0.064)
	0.100	9.399% (4.837%)	8.538 (0.835)	0.110 (0.040)
5.000	0.001	36.115% (7.827%)	12.566 (0.787)	0.261 (0.031)
	0.010	19.423% (0.092%)	15.058 (0.299)	0.325 (0.003)
	0.050	9.832% (5.281%)	8.683 (1.010)	0.125 (0.044)
	0.100	9.304% (4.606%)	8.431 (0.832)	0.110 (0.035)
10.000	0.001	36.115% (7.827%)	12.565 (0.787)	0.261 (0.031)
	0.010	19.421% (0.089%)	14.967 (0.235)	0.325 (0.006)
	0.050	9.822% (5.157%)	8.781 (0.919)	0.135 (0.035)
	0.100	9.758% (5.182%)	8.341 (0.738)	0.105 (0.030)

TABLE 21. IMPACT OF INVERSION ON THE NOISING AND CLIPPING FEDAVG ALGORITHM ON THE SVHN DATASET WITH THE CNN1 MODEL

C	δ	Final ACC \uparrow	PSNR \downarrow	SSIM \downarrow
0.100	0.001	36.413% (2.232%)	13.804 (0.218)	0.296 (0.002)
	0.010	18.947% (0.672%)	13.395 (3.252)	0.278 (0.099)
	0.050	10.065% (0.863%)	14.256 (0.456)	0.300 (0.010)
	0.100	10.758% (2.054%)	14.210 (0.470)	0.299 (0.012)
0.500	0.001	79.505% (3.458%)	14.747 (0.594)	0.317 (0.009)
	0.010	19.481% (0.064%)	13.496 (1.648)	0.283 (0.044)
	0.050	9.699% (0.440%)	14.250 (0.459)	0.301 (0.011)
	0.100	10.727% (1.945%)	14.255 (0.501)	0.301 (0.014)
1.000	0.001	82.510% (2.627%)	15.281 (0.480)	0.312 (0.006)
	0.010	19.565% (0.010%)	13.269 (2.473)	0.272 (0.078)
	0.050	8.527% (0.742%)	14.023 (0.216)	0.293 (0.008)
	0.100	10.748% (1.898%)	14.257 (0.456)	0.304 (0.012)
5.000	0.001	80.592% (2.830%)	15.343 (0.492)	0.317 (0.002)
	0.010	18.365% (0.097%)	13.036 (2.312)	0.261 (0.068)
	0.050	10.177% (0.421%)	13.969 (0.178)	0.292 (0.003)
	0.100	10.966% (1.929%)	14.262 (0.402)	0.305 (0.010)
10.000	0.001	80.705% (2.553%)	15.389 (0.456)	0.317 (0.009)
	0.010	14.365% (6.143%)	13.311 (2.523)	0.270 (0.070)
	0.050	10.630% (4.328%)	14.145 (0.271)	0.296 (0.007)
	0.100	10.306% (0.974%)	14.156 (0.463)	0.302 (0.010)

TABLE 22. IMPACT OF INVERSION ON THE NOISING AND CLIPPING FEDAVG ALGORITHM ON THE SVHN DATASET WITH THE CNN2 MODEL

C	δ	Final ACC \uparrow	PSNR \downarrow	SSIM \downarrow
0.100	0.001	19.590% (0.000%)	15.436 (0.116)	0.322 (0.011)
	0.010	18.730% (1.268%)	14.972 (0.150)	0.317 (0.010)
	0.050	14.738% (4.961%)	15.077 (0.129)	0.317 (0.012)
	0.100	10.365% (4.928%)	15.166 (0.069)	0.318 (0.011)
0.500	0.001	19.590% (0.000%)	15.425 (0.194)	0.321 (0.004)
	0.010	19.538% (0.042%)	14.198 (0.713)	0.299 (0.019)
	0.050	15.433% (5.250%)	15.365 (0.151)	0.319 (0.012)
	0.100	12.275% (3.420%)	15.088 (0.121)	0.317 (0.012)
1.000	0.001	19.590% (0.000%)	15.534 (0.180)	0.323 (0.005)
	0.010	19.577% (0.013%)	14.898 (0.575)	0.312 (0.017)
	0.050	16.189% (5.887%)	15.350 (0.200)	0.319 (0.012)
	0.100	16.346% (2.048%)	15.224 (0.324)	0.318 (0.013)
5.000	0.001	19.590% (0.000%)	15.445 (0.288)	0.316 (0.009)
	0.010	19.588% (0.010%)	14.723 (0.809)	0.312 (0.019)
	0.050	16.742% (4.912%)	15.305 (0.230)	0.319 (0.012)
	0.100	8.245% (1.580%)	15.299 (0.145)	0.319 (0.011)
10.000	0.001	19.590% (0.000%)	15.471 (0.256)	0.316 (0.008)
	0.010	19.595% (0.009%)	14.626 (0.392)	0.302 (0.006)
	0.050	7.763% (1.391%)	15.315 (0.187)	0.319 (0.012)
	0.100	14.898% (5.291%)	15.205 (0.074)	0.318 (0.011)

8. Effectiveness of Our Algorithm Against iDLG

In Table 23 we evaluated the effectiveness of the improved Deep Leakage from Gradients attack [8] under the same settings as our experiments in Section ??, barring batch size which was set to 1 due to the limitation of the attack. We have only evaluated the LeNet and CNN1 models due to memory resource exhaustion when evaluating CNN2 arising from iDLG’s reliance on the LBFGS optimizer, which calculates a second order gradient that is intensive on memory and computation. From our results, generally we see a similar pattern to that from our earlier evaluation with the representation inversion attack, that is, our algorithm achieves greater accuracy² and lower PSNR and SSIM values. The exception in the accuracy performance of the LeNet model trained on SVHN still remains in this case. However, there is a unique exception in this case. Our proposed algorithm experiences larger PSNR and SSIM values on the FMNIST task, that is, it more susceptible to the iDLG attack is in this case. The reason for this occurrence is the combination of the single sample batch size and a simpler dataset, resulting in a running gradient average, m_t , that leads to this resembling more like a single sample.

9. Proof of Statement in Section 5.1

In Section 5.1 we make the following statement,

$$\frac{1}{\sqrt{|\mathcal{C}|^{-1} \sum_t \sum_i g_{t,i}^2}} \leq \frac{1}{\sqrt{\sum_t (|\mathcal{C}|^{-1} \sum_i g_{t,i})^2}}.$$

This can be proven directly as follows.

2. The accuracy values are generally lower due to the single sample minibatch size used within these experiments.

TABLE 23. COMPARISON OF THE IMPACT OF THE iDLG ATTACK AGAINST CANONICAL FEDERATED AVERAGING AND OUR PROPOSED ALGORITHM

Algorithm	Dataset	Model	Accuracy \uparrow	PSNR \downarrow	SSIM \downarrow
FedAvg	FMNIST	LeNet	81.772% (0.908%)	6.309 (0.044)	0.030 (0.017)
		CNN1	84.342% (0.450%)	6.433 (0.145)	0.026 (0.012)
	CIFAR-10	LeNet	27.813% (3.744%)	6.525 (0.214)	0.021 (0.004)
		CNN1	31.093% (7.916%)	8.151 (0.575)	0.037 (0.021)
	SVHN	LeNet	20.751% (2.014%)	7.259 (0.314)	0.007 (0.003)
		CNN1	19.588% (0.000%)	9.945 (0.788)	0.088 (0.006)
Ours	FMNIST	LeNet	82.542% (0.507%)	6.779 (0.085)	0.052 (0.011)
		CNN1	86.232% (0.316%)	6.522 (0.135)	0.035 (0.016)
	CIFAR-10	LeNet	31.913% (0.256%)	6.080 (0.059)	0.009 (0.004)
		CNN1	48.055% (3.461%)	7.602 (0.223)	0.026 (0.003)
	SVHN	LeNet	19.591% (0.004%)	7.598 (0.655)	0.027 (0.009)
		CNN1	26.206% (11.462%)	8.813 (0.480)	0.055 (0.027)

Proof. First we take the reciprocal of both sides,

$$\sqrt{|\mathbb{C}|^{-1} \sum_t \sum_i g_{t,i}^2} \geq \sqrt{\sum_t (|\mathbb{C}|^{-1} \sum_i g_{t,i})^2}.$$

Then we observe the constant, $|\mathbb{C}|^{-1}$, which can be directly drawn out of the surd on the left hand side, and can be first extracted from the parenthesis on the right hand side,

$$\sqrt{|\mathbb{C}|^{-1}} \sqrt{\sum_t \sum_i g_{t,i}^2} \geq \sqrt{\sum_t |\mathbb{C}|^{-2} \sum_i g_{t,i}^2}.$$

The constant, $|\mathbb{C}|^{-2}$, on the right hand side can be drawn out of surd,

$$\sqrt{|\mathbb{C}|^{-1}} \sqrt{\sum_t \sum_i g_{t,i}^2} \geq \sqrt{|\mathbb{C}|^{-2}} \sqrt{\sum_t \sum_i g_{t,i}^2}.$$

Finally, both sides can be divided by the $\sqrt{\sum_t \sum_i g_{t,i}^2}$ term to result in,

$$\sqrt{\frac{1}{|\mathbb{C}|}} \geq \frac{1}{|\mathbb{C}|}. \quad (1)$$

Where (1) holds true since there will always be one or more clients in an operational system. \square

References

- [1] H. Xiao, K. Rasul, and R. Vollgraf, “Fashion-MNIST: a novel image dataset for benchmarking machine learning algorithms,” *arXiv preprint arXiv:1708.07747*, 2017.
- [2] Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner, “Gradient-based learning applied to document recognition,” *Proceedings of the IEEE*, vol. 86, no. 11, pp. 2278–2324, 1998.
- [3] A. Krizhevsky, G. Hinton *et al.*, “Learning multiple layers of features from tiny images,” *Master’s thesis, Department of Computer Science, University of Toronto*, 2009.
- [4] Y. Netzer, T. Wang, A. Coates, A. Bissacco, B. Wu, and A. Y. Ng, “Reading digits in natural images with unsupervised feature learning,” in *NIPS Workshop on Deep Learning and Unsupervised Feature Learning 2011*, 2011. [Online]. Available: http://ufldl.stanford.edu/housenumbers/nips2011_housenumbers.pdf
- [5] I. J. Goodfellow, J. Shlens, and C. Szegedy, “Explaining and harnessing adversarial examples,” *arXiv preprint arXiv:1412.6572*, 2014.
- [6] M. Zaheer, S. J. Reddi, D. Sachan, S. Kale, and S. Kumar, “Adaptive methods for nonconvex optimization,” in *Proceedings of the 32nd International Conference on Neural Information Processing Systems*, ser. NIPS’18. Curran Associates Inc., 2018, p. 9815–9825.
- [7] H. B. McMahan, D. Ramage, K. Talwar, and L. Zhang, “Learning differentially private recurrent language models,” in *International Conference on Learning Representations*, 2018. [Online]. Available: <https://arxiv.org/abs/1710.06963>
- [8] B. Zhao, K. R. Mopuri, and H. Bilen, “iDLG: Improved deep leakage from gradients,” *arXiv preprint arXiv:2001.02610*, 2020.